Fast Parallel SAME Gibbs Sampling on General Discrete Bayesian Networks

Daniel Seita, Haoyu Chen & John Canny

seita@berkeley.edu, haoyuchen@berkeley.edu, jfc@cs.berkeley.edu

Introduction

A fundamental task in machine learning and related fields is to perform inference on Bayesian networks. Since exact inference takes exponential time, it is common to use an approximate algorithm such as Gibbs sampling, but this can still be intractable for graphical models with just a few hundred binary random variables. In this project, we:

- Build a highly optimized Gibbs sampler

- Apply SAME to reduce variance

- Benchmark our Gibbs sampler on "MOOC".

Fast Parallel SAME Sampling

SAME (State Augmentation for Marginal Estimation) [1, 2, 3] can be viewed as cooling the posterior parameter distribution and allows annealed search for the MAP parameters, often yielding very high quality (lower loss) estimates. Given a distribution $P(X, Z \mid \Theta)$, to estimate the most likely Θ based on the data (X, Z) using SAME, one would define a new joint Q:

$$Q(X, \Theta, Z^{(1)}, \dots, Z^{(m)}) = \prod_{j=1}^{m} P(X, \Theta, Z^{(j)})$$

which models m copies of the distribution tied to the same set of parameters Θ.

In order to parallel the sampling, we apply graph coloring to the moralized graph of the original network, and within each color group, sample all the variables in parallel. Left figure shows a simple example of three color groups with m=2.

CPT
node in color group 1
node in color group 2

Figure 1: Gibbs sampler framework (m=2)

Conclusion

We conclude that our Gibbs sampler is much faster than the state of the art (JAGS) in Gibbs sampling and can be applied to data with hundreds of variables. We also argue that SAME is beneficial for Gibbs sampling, and that it should be the go-to method for researchers who wish to perform inference on (discrete) Bayesian networks. Future work will explore the application of our sampler to a wider class of real-world datasets.

References

[1] A. Doucet, S. Godsill, and C. Robert. Marginal maximum a posteriori estimation using markov chain monte carlo. Statistics and Computing, 12:77-84, 2002.

[2] C. Robert, A. Doucet, and S. Godsill. Marginal MAP estimation using markov chain monte-carlo. In IEEE Int. Conf. on Acoustics, Speech and Signal Processing, volume 3, pages 1753-1756. IEEE, 1999.

[3] Z. Huasha, J. Biye, and C. John. Same but different: Fast and high quality gibbs parameter estimation. In Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, KDD '15, pages 1495-1502, NY, USA, 2015. ACM

Implementation

Figure 2: Gibbs sampler framework

Our Gibbs sampler is implemented and integrated as part of the open-source BID-Mach library for machine learning. Figure 2 shows a visualization of how it works on real data. Our sampler expects a (usually sparse) data matrix, with rows representing variables and columns representing cases. BIDMach divides data into same-sized "mini-batches" and iterates through them to update parameters. Going through all mini-batches is one full pass over the data.

Experiments

We benchmark our code on a nation-wide examination dataset, which contains the assessment (correct or not) of student responses to questions. There were 4367 students and 319 questions. Each question is considered an "observed" node in the Bayesian network. We only know 2.2% of the data. We call the student responses data "MOOC" data. We compare our sampler with JAGS.
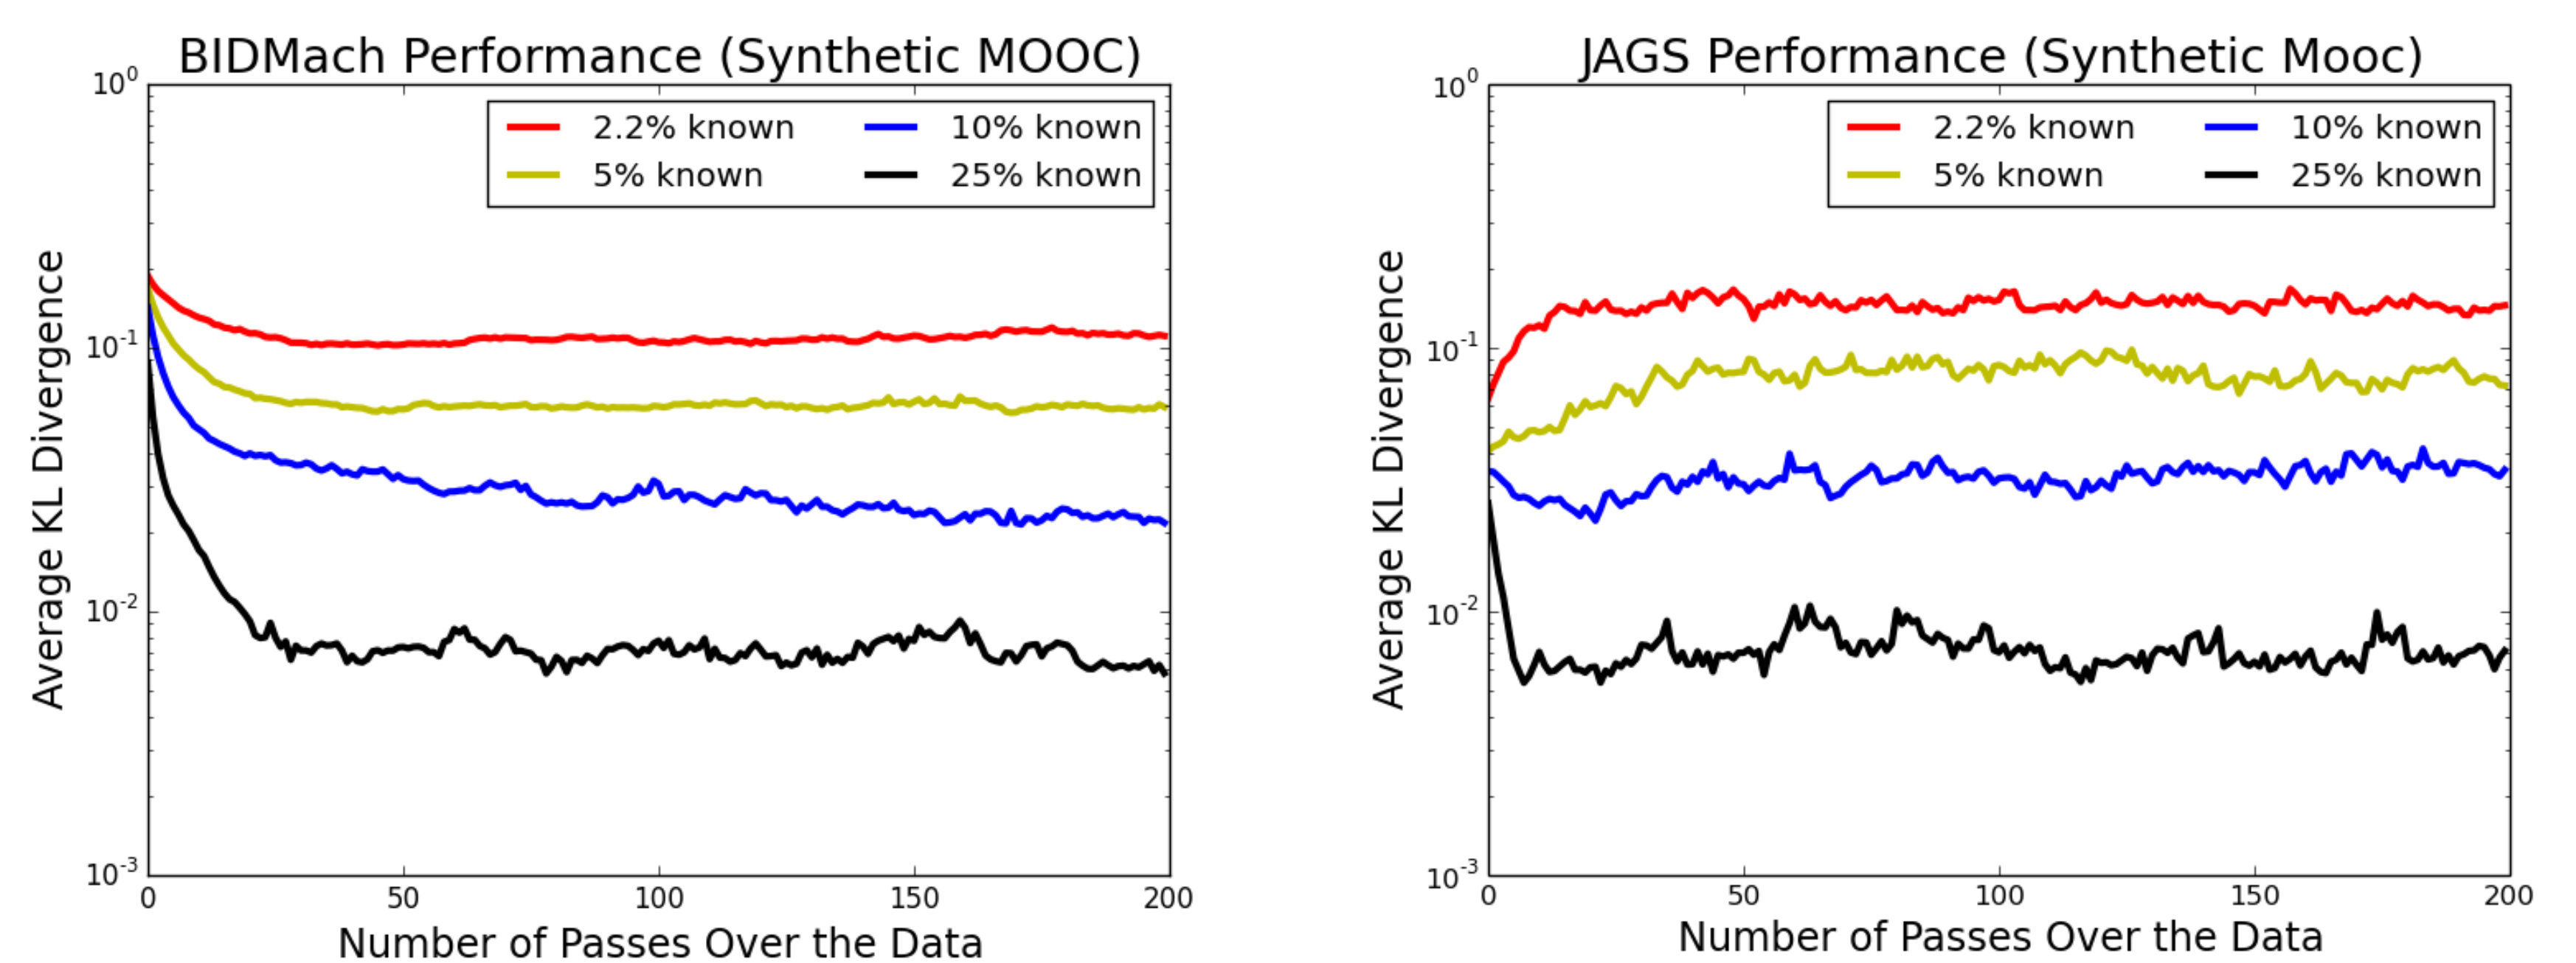

**Figure 3:** The KL_{avg} from BIDMach. (left); The KL_{avg} from JAGS.(right)

Figure 4: The KL_{avg} vs. runtime for BIDMach. (left); Prediction accuracy of BIDMach(right)

In order to show the performance of our sampler in the limit, we replicate the data to increase its size. Table 1 records the runtime (CPU) per iteration for both BIDMach and JAGS, which reveals that our sampler has much shorter runtime than JAGS even without using GPU.

	1x	2x	5x	10x	20x	40x
BIDMach Time(sec)/Iter	0.182	0.375	0.931	1.792	3.501	7.181
JAGS Time(sec)/Iter	4.876	13.745	29.150	94.075	171.545	OOM

Table 1: BIDMach (CPU) vs. JAGS Runtime on (Replicated) Real Data

Table 2 records the BIDMach runtime (GPU) and GigaFlops, which indicates the time and gflops tradeoff.

	$m = 1$	$m = 5$	$m = 10$	$m = 20$	$m = 30$	$m = 40$	$m = 50$
GigaFlops (K)	2.38	6.98	9.08	11.00	11.70	12.27	12.35
Time/Iter (K)	0.307	0.523	0.804	1.328	1.872	2.380	2.957
GigaFlops (RM)	1.65	4.58	7.08	9.69	10.92	11.62	12.02
Time/Iter (RM)	0.666	1.192	1.541	2.252	2.998	3.753	4.536

Table 2: BIDMach (GPU) Runtime vs. GigaFlops on Large Data